Custom Search	Search

RE: about the time consuming kvm_vcpu_ioctl issue in rhel kernel

[Date Prev][Date Next][Thread Prev][Thread Next][Date Index][Thread Index]

- *Subject*: RE: about the time consuming kvm_vcpu_ioctl issue in rhel kernel
- *From*: "Li, Liang Z" <liang.z.li@xxxxxxxxx>
- *Date*: Thu, 6 Aug 2015 05:58:36 +0000
- *Accept-language*: zh-CN, en-US
- *Cc*: "Zhang, Yang Z" <yang.z.zhang@xxxxxxxxx>, "kvm@xxxxxxxxxxxxxxx" <kvm@xxxxxxxxxxxxxxx>
- *Thread-index*: AdDP9UJwwpuw8Ol+QP2Z1DthOfo0awAE1k6wAADmSiA=
- *Thread-topic*: about the time consuming kvm_vcpu_ioctl issue in rhel kernel

2

```
Please Ignore the following message  in the brace
{ The synchronize_rcu() is a time consuming operation,  the unpstream kernel
 still have some issue, the KVM_RUN ioctl will take more then 10ms when
 resume the VM after migration. }

The upstream kernel does not have such issue, only the rhel kernel has. Forgive my neglect.

Liang


>
>
> > -----Original Message-----
> > From: Li, Liang Z
> > Sent: Thursday, August 06, 2015 11:47 AM
> > To: 'Paolo Bonzini'; 'Juan Quintela'; linux-kernel@xxxxxxxxxxxxxxxx
> > Cc: Zhang, Yang Z; kvm@xxxxxxxxxxxxxxx
> > Subject: about the time consuming kvm_vcpu_ioctl issue in rhel kernel
> >
> > Hi Paolo & Juan,
> >
> > I found some of the kvm_vcpu_ioctl operation takes about more than
> > 10ms with the 3.10.0-229.el7.x86_64 kernel, which prolong the VM
> > service downtime when doing live migration about 20~30ms.
> > This happened when doing the KVM_KVMCLOCK_CTRL ioctl. It's worse if
> > more VCPUs used by guest(4VPUs in my test). The upstream kernel does
> > not have this issue because of the patch 7a72f7a140bfd3,  maybe this
> > patch should be backported to the el7 kernel.
> >
```

```
> > Liang
--
To unsubscribe from this list: send the line "unsubscribe kvm" in
the body of a message to majordomo@xxxxxxxxxxxxxxx
More majordomo info at   http://vger.kernel.org/majordomo-info.html
```

- Prev by Date: **RE: about the time consuming kvm_vcpu_ioctl issue in rhel kernel**
- Next by Date: **RE: [PATCH 5/9] KVM: x86: unify handling of interrupt window**
- Previous by thread: **RE: about the time consuming kvm_vcpu_ioctl issue in rhel kernel**
- Next by thread: **apparent SMBASE relocation issue with noexec enabled [was: MdeModulePkg DxeIpl: Add stack NX support]**
- Index(es):
    - **Date**
    - **Thread**

[Index of Archives]   [KVM ARM]   [KVM ia64]   [KVM ppc]   [Virtualization Tools]   [Spice Development]   [Libvirt]   [Libvirt Users]   [Linux USB Devel]   [Linux Audio Users]   [Yosemite Questions]   [Linux Kernel]   [Linux SCSI]   [XFree86]